

Text Analytics – A Comparative Analysis of Brand Reviews

Jaime Rubio Diaz

40425150

Contents

1. Introduction and Background	3
2. Methodology	3
3. Text Data Pre-Processing	7
4. Text Analytic	8
4.a. Supervised & Semi-Supervised Machine Learning	8
4.b. Visual analytics.....	11
4. Results and Discussion.....	12
5. Conclusion and Recommendations	15
6. References	17
7. Python Code	20
8. Appendix.....	39
Table 1. Data description	4
Table 2. Supervised learning models	6
Table 3. “TfidfVectorizer” parameters selection	8
Table 4. Supervised learning models	9
Table 5. Hyperparameter explanation for the best performance model.....	9
Table 6. Potential reasons for performance and/or limitations supervised learning.....	9
Table 7. Semi-supervised learning models	10
Table 8. Potential reasons for performance and/or limitations SSL	10
Table 9. Appropriate clasification metrics	12
Figure 1. CRISP-DM	3
Figure 2. First data visualizations	4
Figure 3. Cleaning text process.....	5
Figure 4. Reviews language distribution	7
Figure 5. Reviews language distribution after the translation	7
Figure 6. Results obtained from test pre-processing.....	8
Figure 7. Predicted labels after filling missing values with predictions	11
Figure 8. Emotions distributions for each brand.....	11
Figure 9. Start rating by emotion for Z_.....	11
Figure 10. Start rating by emotion for H_.....	12
Figure 11. Accuracy comparison supervised models	12
Figure 12. SL SDG classification report and confusion matrix.....	13
Figure 13. SSL SDG classification report and confusion matrix.....	13

1. Introduction and Background

The objective of this report is to outline the analysis conducted on customer reviews of two brands, using a dataset of 5,722 entries. Given the significant number of missing entries in the "emotions" column, a semi-supervised learning (SSL) approach is adopted to accurately label these missing observations, drawing on both the existing labels and the customer reviews. This method enables an analysis of emotional responses across both brands. The insights gained from this analysis provide the brands with a better understanding of their performance.

Numerous studies focus on the sentiment analysis of customer reviews, due to the valuable insights they offer companies. Singla, et. al., 2017 using a data set of 400,000 reviews, carried out a supervised learning analysis to develop a classifier that categorizes reviews into two sentiment types: positive and negative. These approaches are similar to those initially employed in this project, utilizing the labelled data.

Regarding semi-supervised learning, in the context of detecting fake opinions in hotel customer reviews, Ligthart et al., (2021) found that self-training algorithms based on Naïve Bayes, maintained the performance of a supervised model with only a small portion of the data labelled. Conversely, Widmann and Verberne (2017), explored graph-based algorithms like label propagation, using the 20-Newsgroup corpus dataset, highlighting the high interpretability of these algorithms in a semi-supervised context, particularly for identifying key phrases in documents or understanding their relationships to labels.

2. Methodology




Figure 1. CRISP-DM
(Source: Wirth and Hipp, 2000)

Business Understanding

The nature of the business is unknown, but the presence of reviews in over 20 languages and mentions of delivery issues in those reviews suggest that both brands operate online. Reputation is the most important thing for online stores, as customers frequently view the reviews of others before making a purchase, therefore analysing and addressing issues highlighted in these reviews can lead to more positive reviews and attract more customers (Halim et al., 2022).

Data Understanding

Table 1. Data description

Variables		
ID_	Unique number for each customer	Categorical
brand_name_	Name of the brand	Categorical
country_	Customer country	Categorical
star_rating_	Rating from 1 to 5	Numerical (Ordinal)
emotions_	Emotion of the customer review	Categorical
text_reviews_	Text review of the customers	text

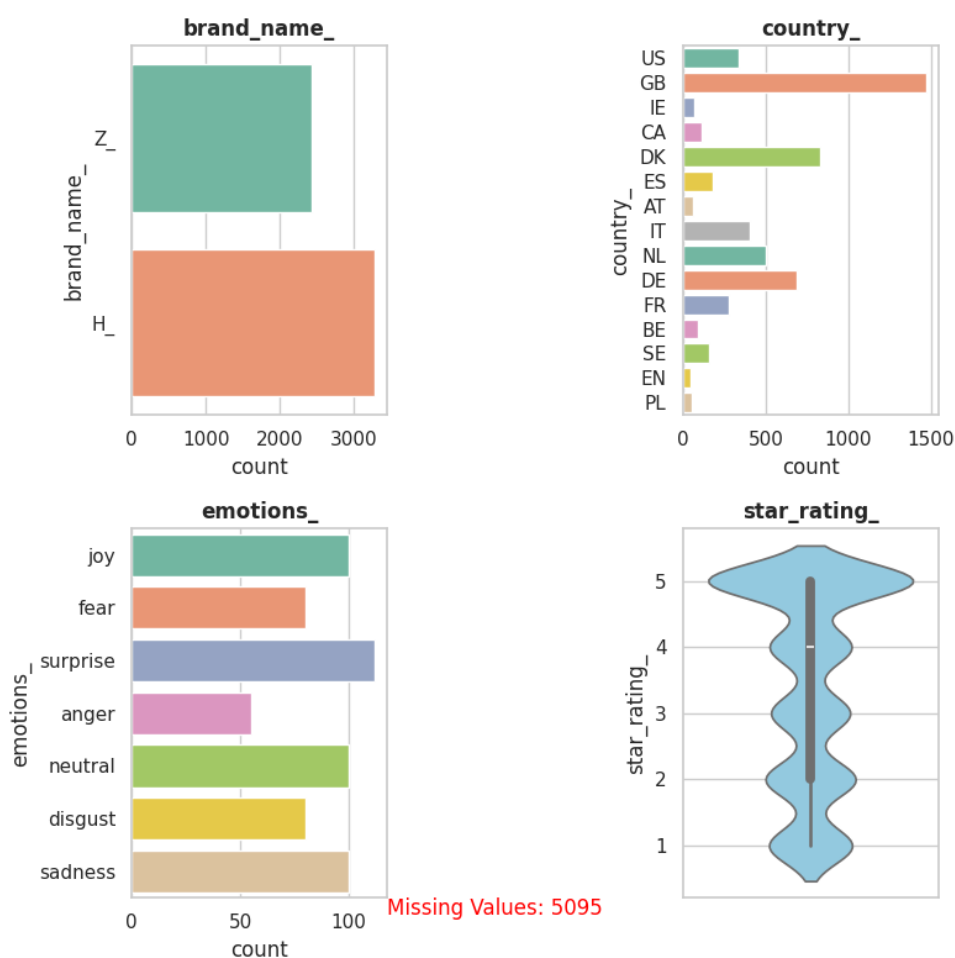


Figure 2. First data visualizations

(** Given the extensive variety of countries represented in the dataset, only the top 15 are shown in the graph.)

The dataset stores reviews for 2 brands, namely “Z_” and “H_”. Notably, the dataset reveals that the predominant rating by customers is 5 stars. In terms, of emotions “surprise” is the most commonly reported. However, it is important to highlight that 5095 reviews do not have specified emotional responses. Furthermore, the data indicates that Great Britain (GB) is the most frequently mentioned country. Given that all the data is categorical, except for the star rating variable which is ordinal, frequency-related graphs such as bar charts or pie charts are suitable for the categorical data. For the

distribution of the categories and the ordinal data, violin plots can be used. Violin plots are particularly useful compared to boxplots because they not only display the median and interquartile ranges but also provide a detailed visualization of the entire data distribution, including potential multimodality (Hintze & Nelson, 1998).

Data Preparation

The variables “brand_name_”, “country_”, “star_rating_”, and “emotions_” do not require pre-processing. However, “text_reviews_” variable needs to be pre-processed. As shown in [Figure 3](#), the reviews are in various languages, but they need to be translated into English for accurate prediction and analysis of the unlabelled emotions. For this project, Deepl API and Baidu API are used for translation. The Deepl API allows for the translation of up to 500,000 text characters for free, hence Baidu API is used for the remaining translations. After translation, a new data set is generated where all the reviews are in English.

Secondly, to manipulate text in Python it is necessary to clean it. Cleaning the text can enhance the performance of the algorithm by removing noise and reducing the complexity of the text. The following graph illustrates the steps taken in the data cleaning process.

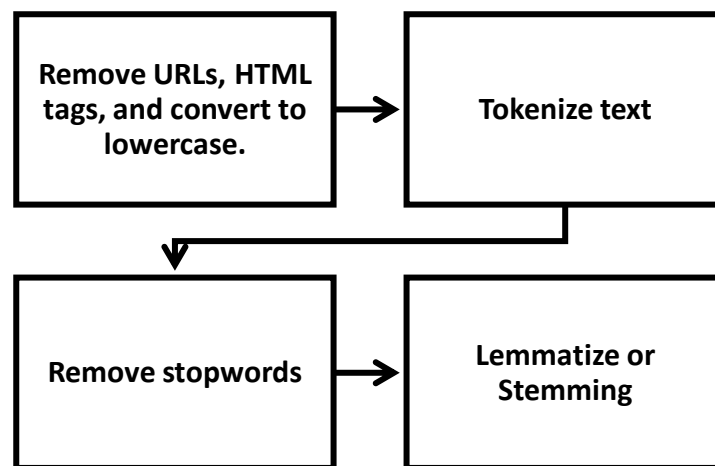


Figure 3. Cleaning text process

The process includes the removal of URLs, HTML tags and the conversion to lowercase, as well as the removal of stopwords, which are elements or words that do not contribute meaningfully to the document, after this the text is tokenized, converting it into individual words (Vijayarani et al., 2015). Lemmatization and stemming are both text normalization techniques aimed at word reduction. Stemming simplifies words by removing endings through basic suffix elimination, often resulting in non-words; while lemmatization is a more time-complex process that considers the part of speech of a word and uses a lexical resource to ensure that the word reduction corresponds to a valid lemma (Chai, 2022).

Once the text is cleaned, it needs to be converted into a numerical format that machine learning algorithms can process. Two of the most common techniques are Bag of Words (BoW) and Term Frequency-Inverse Document Frequency (TF-IDF). BoW focuses on the presence or absence of words in a document, while TF-IDF takes into account not only the presence, but also the relative importance of the words within the sentence in relation to the document (Pimpalkar and Retna, 2020).

$$TF - IDF = \frac{(\text{Times term appears in the sentence})}{(\text{Total terms in the sentence})} * \log \frac{(\text{Total sentences})}{(\text{Total sentences with the term})}$$

(Source: Pimpalkar and Retna, 2020)

Although using all text features can result in good performance, feature selection often enhances classifier performance (Bird et al., 2009). When using “tfidfvectorizer”, a parameter can be set to select the number of features for each review, choosing the most frequent words for that review. Additionally, “ngram_range” parameter can be customized during vectorization, determining the range of word combinations considered. Both parameters are evaluated with a Logistic Regression model.

Modelling

In a semi-supervised learning context, an initial supervised phase is conducted to leverage labelled data and to establish a baseline mode. For text classification with a target variable of five levels, Gurcan (2018) concluded that Multinomial Naïve Bayes algorithm, with an accuracy of 90%, was the most effective for the task. Furthermore, Joshi and Abdelfattah (2021) conducted research to predict medical conditions associated with drug effects based on online pharmaceutical reviews using six different supervised machine learning classifiers, finding that the Linear Support Vector Classifier performed the best. Based on these findings both algorithms will be tested with our data, in combination with other classifiers.

Table 2. Supervised learning models

Supervised Learning Classifiers	
Random Forest	Random Forest expands upon the decision tree model by generating a multitude of trees in a randomized way, enhancing accuracy while preserving interpretability (Rigatti, 2017). Despite the robustness, irrelevant features can distract the model.
K-Nearest Neighbours (KNN)	The K-Nearest Neighbours classifier assigns a category to an observation by analysing the categories of the nearest examples within the dataset, it calculates proximity using distance metrics like Euclidean or Manhattan distance to identify similarities (Zhang, 2016). The performance of KNN decreases with high-dimensional data.
Support Vector Machine	This algorithm employs a nonlinear mapping to transform the data to form a hyperplane in a transformed space, where it performs classification using kernel functions (Hearst et al., 2002). SVM requires careful tuning of hyperparameters, and the selection of an appropriate kernel function.
Stochastic Gradient Descent (SGD)	Stochastic Gradient Descent (SGD) is an optimization technique, particularly effective for handling large-scale and sparse data sets, that can employ various loss functions to train models efficiently (<i>scikit, no date</i>).
Multinomial Naïve Bayes	MNB is a variant of the Naïve Bayes algorithm, often used in text classification, which is particularly useful for feature vectors where the attributes represent the frequencies by a multinomial distribution (McCallum and Nigam, 1998).

Once supervised approach is performed, a semi-supervised approach is employed for labelling the unlabelled data. Semi-supervised learning, a machine learning approach that combines labelled data with a large amount of unlabelled data, is particularly useful as labelling can be time-consuming and expensive (Zhu, 2005). Van Engelen and Hoos (2020), categorize semi-supervised techniques into inductive methods, which use all the data to create a model that can predict labels for new data, and transductive methods, that are focused on the propagation of the labels based on direct connections. Common algorithms in the first category include wrapper methods such as Self-training, while in the

second category Graph-based methods. Graph-based methods propagate labels by using a graph $G = (V, E)$, where V represents pairs of labelled and unlabelled data, and E a matrix indicating similarity between V pairs, while Self-training, generates pseudo-labels to expand the training data set with labels that enhance the performance, resulting in a model similar to supervised approach (Duarte and Berton, 2023).

After semi-supervised, the best-performing model is selected, and it is used to predict the unlabelled data. These predictions are incorporated into the original dataset, resulting in a dataset without missing values, allowing for an analysis of the brand reviews.

For the analysis, considering the abundance of text features, a logistic regression model can be constructed to examine the importance of coefficients associated with specific words (Müller and Guido, 2016 pag. 426-431). The coefficients from the model are then plotted to identify the most influential words for predicting each emotion. This approach provides an overview of key linguistic indicators, revealing the most impactful words in customer reviews associated with each emotion.

3. Text Data Pre-Processing

As detailed above, the first step was to translate the text reviews into English. Figure 3 illustrates the distribution of languages in the original reviews, and Figure 4 shows the distribution after translation. Although other languages still appear in Figure 4, they are present in a negligible proportion. This may be because the API failed to recognize these languages, and therefore, did not translate them, or because the "langdetect" library incorrectly identified them due to the presence of special characters in those observations. Regardless, this is not a concern due to the small amount they represent and the overall improvement over the original reviews. Now, 5,583 reviews are in English, compared to only 2,831 in the original dataset.




Figure 4. Reviews language distribution

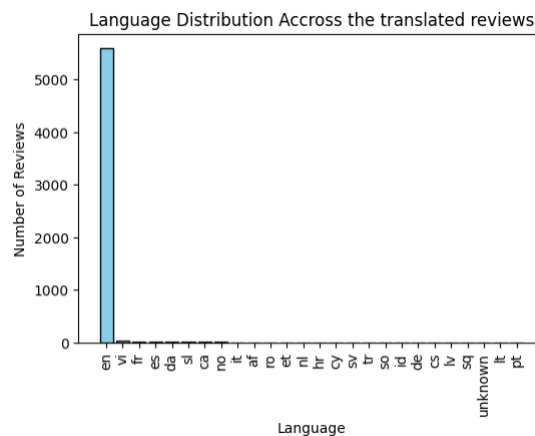


Figure 5. Reviews language distribution after the translation

The methodology section detailed the steps followed for cleaning the text. Here, the results after each step are shown for the first review:

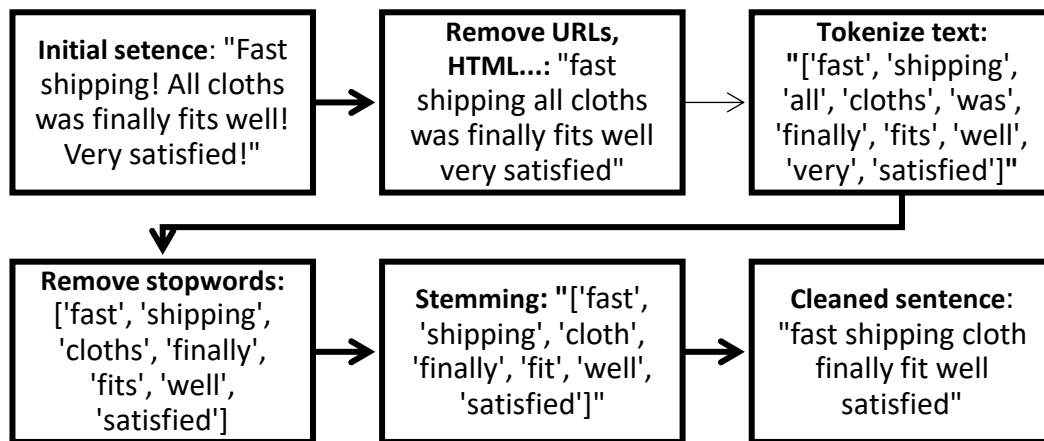


Figure 6. Results obtained from test pre-processing

From the initial review text, "Fast shipping! All cloths was finally fits well! Very satisfied!", the text cleaning and normalization process transforms it into "fast shipping cloth finally fit well satisfied". In this example, stemming was chosen to normalize the text. This decision was based on experimentation within a logistic regression model, where stemming led to better performance.

After cleaning, the text is vectorized using the TfidfVectorizer. Different values for the ngram_range and max_features parameters are tested within a Logistic Regression model to find the most effective combination for handling the text data.




Table 3. "TfidfVectorizer" parameters selection

TfidfVectorizer selected parameters	
n_features	3000
ngram	(1,2)

4. Text Analytic

4.a. Supervised & Semi-Supervised Machine Learning

Initially, the data set is split into labelled data and unlabelled data. For the supervised learning models the labelled data is split into train data set (80%) and test data set (20%). Once split, the supervised models are constructed and their performance, in terms of accuracy, is assessed after hyperparameter tuning.

Table 4. Supervised learning models

Classifier	Best Hyperparameters	Accuracy
Random Forest	{'clf__max_depth': 20, 'clf__min_samples_split': 10, 'clf__n_estimators': 200}	0.60
K-Nearest Neighbours	{'clf__algorithm': 'auto', 'clf__n_neighbors': 7, 'clf__weights': 'distance'}	0.36
Support Vector Machine	{'clf__C': 1, 'clf__kernel': 'linear'}	0.46
SDG	{'clf__alpha': 0.001, 'clf__loss': log, 'clf__penalty': 'l1'}	0.61
Multinomial Naïve Bayes	{'clf__alpha': 0.01}	0.47

The highest accuracy among supervised classifiers is shown by Stochastic Gradient Descent (SDG) with the loss function set to “log”.

Table 5. Hyperparameter explanation for the best performance model

Source: scikit (no date)

Hyperparameter explanation for SDG	
clf_loss	It is the loss function used by the model. ‘log’ = logistic regression
penalty	Regularization applied. ‘l1’ = commonly referred as lasso.
alpha	Used to compute learning rate and multiples the regularization.

Table 6. Potential reasons for performance and/or limitations supervised learning.

Classifier	Potential Reasons for Performance and/or Limitations
Random Forest	It demonstrates strong performance, probably stemming from its capability to effectively manage non-linear and categorical relationships.
K-Nearest Neighbours	The distance metrics for KNN performs worse in high-dimensional spaces, which is the case of text-data (Kouroukidis and Evangelidis, 2011).
Support Vector Machine	Since it requires careful hyperparameter tuning, the small size of the dataset may have influenced the improper selection of hyperparameters.
SDG	SGD is robust to a large number of features in text data, the use of logistic regression (loss=log) and L1 penalty might have helped in text feature selection.
Multinomial Naïve Bayes	The use of td-idf instead of feature counts may have affected the performance of the algorithm (<i>scikit, no date</i>).

For the semi-supervised learning phase, train and test datasets need to be built. The test dataset used for the semi-supervised remains the same as that used in the supervised phase to ensure that the performance of the semi-supervised models is evaluated against labelled data. The train dataset is constructed by combining the training data set used in the supervised models and all unlabelled data.

This results in a test dataset consisting only of labelled data and a train dataset that includes both labelled and unlabelled data.

Following dataset preparation, emotions are mapped and converted to a numerical format in both the train and test datasets. With the data ready, the semi-supervised models are then constructed, employing both “Self-Training” and “Label Spreading”. For the Self-Training approach, the parameters selected are those which provided the best performance in the previous phase.

Table 7. Semi-supervised learning models

Label Spreading	Accuracy
Label Spreading	0.32

SelfTraining Classifier	Accuracy
Random Forest	0.62
Support Vector Machine	0.55
SDG	0.63
Multinomial Naïve Bayes	0.43

Self-training classifier shows better performance than Label Spreading, and just as in the supervised learning phase, the Stochastic Gradient Descent (SGD) classifier demonstrates the best performance among the Self-learning models.

Table 8. Potential reasons for performance and/or limitations SSL

SSL method	Potential Reasons for Performance and/or Limitations
Label Spreading	Label Spreading assumes that similar data points are likely to share labels, which may not accurately reflect the actual relationships in this dataset.
SelfTraining	The main limitation could be that the model may propagate errors if the initial model predictions are inaccurate, leading to amplified errors in subsequent iterations. Duarte and Berton, (2023) discusses some strategies to select more confidently labelled instances.

Utilizing the Self-training SGD classifier, the unlabelled data is effectively labelled, expanding the dataset's utility for further analysis. Following this, in line with the methodology outlined earlier, two logistic regression models are constructed, one for each brand, using the text reviews and emotions columns, obtaining an accuracy of 0.69 for Brand Z_ and 0.75 for Brand H_ (results of the models will be detailed in the “Results and Discussion” section).

4.b. Visual analytics

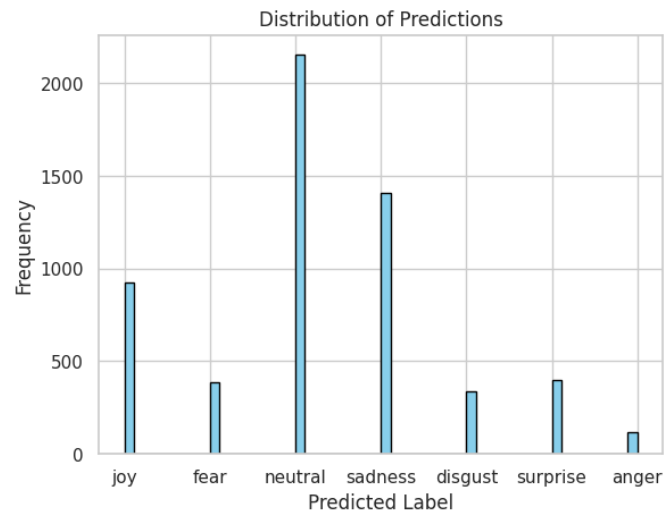


Figure 7. Predicted labels after filling missing values with predictions




Figure 8. Emotions distributions for each brand

Bar charts illustrates a comparative analysis of the emotional responses for the two brands, based on frequency and percentage. For both brands, "neutral" is the predominant emotion, which might suggest either a general sense of impartiality among customers or potentially poor model performance in predicting nuanced emotions. Following neutral, "sadness" is the second most frequent emotion, comprising 25.55% of the reviews for Brand H_ and 23.25% for Brand Z_.




Figure 9. Start rating by emotion for Z_

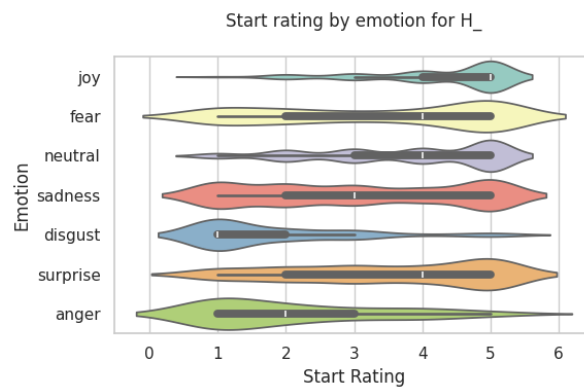


Figure 10. Start rating by emotion for H_

Overall, most of the reviews for both brands have a 5-star rating. Predictions suggest that higher star ratings are commonly associated with emotions such as joy, neutral, and surprise. This indicates that reviews labelled with surprise are generally positive. Conversely, the presence of negative emotions like fear or sadness, even in reviews with a 5-star rating, suggests potential areas for improvement in the model's accuracy.

4. Results and Discussion

For the supervised classifiers k-fold cross validation with k= 10 is performed and the following are the comparison of accuracies:




Figure 11. Accuracy comparison supervised models

SDG has the best performance in term of accuracy so it is further explained. To explain it some metrics will be consider.

Table 9. Appropriated classification metrics

Appropriate metrics	
Confusion matrix	It compares true label with predicted labels.
Precision	It refers to the percentage of correctly identified positive results out of all results classified as positive.
Recall	It measures the proportion of positive patterns that are correctly identified and classified as such.

F1-score	It is the average between precision and recall.
----------	---

(Hossin and Sulaiman, 2015)

Classification Report:					Confusion Matrix							
	precision	recall	f1-score	support		0	1	2	3	4	5	6
anger	0.67	0.36	0.47	11	True Label	4	2	0	2	2	0	1
disgust	0.54	0.44	0.48	16		1	7	0	1	3	2	2
fear	0.85	0.69	0.76	16		0	1	11	0	0	3	1
joy	0.69	0.90	0.78	20		0	0	0	18	0	1	1
neutral	0.56	0.70	0.62	20		1	1	0	0	14	0	4
sadness	0.50	0.35	0.41	20		0	2	1	1	5	7	4
surprise	0.55	0.70	0.62	23		0	0	1	4	1	1	16
accuracy			0.61	126		0	1	2	3	4	5	6
macro avg	0.62	0.59	0.59	126								
weighted avg	0.61	0.61	0.60	126								

Predicted Label

Figure 12. SL SDG classification report and confusion matrix

Overall, the model achieves an accuracy of 0.61. It performs best in identifying 'joy', with a precision of 0.69 and a recall of 0.90. This suggests that while the model is quite effective at detecting instances of 'joy' when they occur, it sometimes misclassifies other emotions as 'joy'. Similarly, it detects 'fear' effectively. Conversely, the model struggles with 'sadness', achieving a precision of 0.50 and a recall of 0.35.

Classification Report:					Confusion Matrix							
	precision	recall	f1-score	support		0	1	2	3	4	5	6
0	1.00	0.36	0.53	11	True Label	4	1	0	2	2	1	1
1	0.80	0.50	0.62	16		0	8	0	1	4	3	0
2	0.67	0.62	0.65	16		0	0	10	0	4	1	1
3	0.66	0.95	0.78	20		0	0	0	19	0	1	0
4	0.48	0.70	0.57	20		0	0	0	1	14	3	2
5	0.50	0.55	0.52	20		0	1	1	3	4	11	0
6	0.76	0.57	0.65	23		0	1	2	3	4	1	12
accuracy			0.63	126								
macro avg	0.70	0.61	0.62	126								
weighted avg	0.67	0.63	0.62	126								

Predicted Label


Figure 13. SSL SDG classification report and confusion matrix

Same that for supervised learning, class 3 (joy) demonstrates the highest performance, with a precision of 0.66 and a recall of 0.95, suggesting that the model is particularly effective in identifying this class. In contrast, Class 0 (anger) demonstrates perfect precision at 1.00 but a low recall of 0.36. This indicates that while all predictions made for Class 0 are correct, a significant number of actual Class 0 instances are not predicted by the model. This is confirmed by the confusion matrix. The overall accuracy of the model is 0.63, indicating a moderate performance. However, considering that there are 7 classes, the random chance of correctly predicting a class is approximately 0.14 (1/7). Moreover, the SSL with a low portion of labelled data enhances the performance of the supervised.

In terms of the logistic regression models, the most recent version of scikit-learn uses the 'lbfgs' solver by default for multiclass classification problems (scikit, no date). 'lbfgs' uses a reference class for comparison with the rest classes, where positive coefficients increase the probability of a given class relative to this reference class. Therefore, in the analysis, it is particularly important to focus on the

positive coefficients. These coefficients highlight words that are significant in predicting a particular emotion.


-. Brand Z_:



**Only these two levels are shown in this section. To increase report quality rest of the graph are shown in the Appendix.

Regarding, Brand Z_, words like “staff” and “rude staff” are strongly associated with the predictions of anger, disgust and sadness. Additionally, terms such as "worst customer" and "customer service" frequently appear in disgust, sadness and fear. Moreover, expressions related to returns are highly significant for fear and anger emotions. In the case of sadness, which is the second most predicted emotion, there are words such as “slow” or “slow delivery”.

-. Brand H_:



For Brand H₂ in the context of negative emotions: for anger, words such as "rude," "angry," and "lazy" appear, disgust is influenced by harsh descriptors like "dirty" or "poor" probably indicating dissatisfaction with aspects related to product quality and other terms such as "worst customer", "service even" or "staff". Sadness, similar as brand Z₂, is marked by words such as by "service", "staff" and "slow," illustrating scenarios where services or deliveries fell below customer expectations. In the context of positive emotions, for joy emotion "price good" accounts as a good review.

5. Conclusion and Recommendations

Based on the analysis some actionable insights and recommendation can be given to the brands. Firstly, for brand Z₂ it is evident that there is substantial room for improvement in the areas of customer service, delivery processes, and return mechanisms. Secondly, Brand Z₂ needs to consider an improvement in customer service, delivery processes and product quality, while the currently products' price seems to be valued by the customers.

Customer service can be improved by the development of training programs for customer service teams to improve interaction quality and response times. For delivery process, logistics can be analysed to detect how to reduce delivery times. Moreover, it is recommended to simplify the return process by making it easier. These are general recommendations, if more information needs to be extracted it can be created a way to ask customers about these processes to obtain more inputs from them.

For future work, it is suggested to explore other semi-supervised techniques that could potentially provide better insights and performance. For instance, advanced neural network models may offer better performance. Using a convolutional neural network with a ladder network in the context of semi-supervised learning (SSL) and image classification, Liu et al. (2017) achieved an accuracy of 97.6% when labelling the unlabelled data.

6. References

- 1) 1.5. *stochastic gradient descent* (no date) *scikit*. Available at: <https://scikit-learn.org/stable/modules/sgd.html> (Accessed: 04 April 2024).
- 2) Bird, S., Klein, E. and Loper, E. (2009) *Natural language processing with python*. Beijing: O'Reilly.
- 3) Chai, C.P. (2022) 'Comparison of text preprocessing methods', *Natural Language Engineering*, 29(3), pp. 509–553. doi:10.1017/s1351324922000213. Available at: <https://www.cambridge.org/core/journals/natural-language-engineering/article/comparison-of-text-preprocessing-methods/43A20821D65F1C0C4366B126FC794AE3> (Accessed: 04 April 2024)
- 4) Duarte, J.M. and Berton, L. (2023) 'A review of semi-supervised learning for text classification', *Artificial Intelligence Review*, 56(9), pp. 9401–9469. doi:10.1007/s10462-023-10393-8. Available at: <https://link.springer.com/article/10.1007/s10462-023-10393-8> (Accessed: 20 April 2024)
- 5) Gurcan, F. (2018) 'Multi-class classification of Turkish texts with machine learning algorithms', 2018 2nd International Symposium on Multidisciplinary Studies and Innovative Technologies (ISMSIT) [Preprint]. doi:10.1109/ismsit.2018.8567307. Available at: <https://ieeexplore.ieee.org/abstract/document/8567307> (Accessed: 10 April 2024)
- 6) Halim, E., Wardhana, G. and Sasongko, A.H. (2022) 'Online customer reviews as a marketing tool to generate customer purchase intention in Ecommerce', 2022 IEEE International Conference on Cybernetics and Computational Intelligence (CyberneticsCom) [Preprint]. doi:10.1109/cyberneticscom55287.2022.9865598. Available at: <https://ieeexplore.ieee.org/abstract/document/9865598> Accessed: 04 April 2024
- 7) Hearst, M.A. et al. (2002) 'Support Vector Machines', *IEEE Xplore*, 13(4), pp. 18–28. doi:10.1142/9789812776655_0002. Available at: <https://ieeexplore.ieee.org/document/708428/citations#citations> Accessed: 10 April 2024
- 8) Hintze, J.L. and Nelson, R.D. (1998) 'Violin plots: A box plot-density trace synergism', *The American Statistician*, 52(2), p. 181. doi:10.2307/2685478. Available at: <https://www.tandfonline.com/doi/abs/10.1080/00031305.1998.10480559> Accessed: 20 April 2024.
- 9) Hossin, M. and Sulaiman, M.N. (2015) 'A review on evaluation metrics for data classification evaluations', *International Journal of Data Mining & Knowledge Management Process*, 5(2), pp. 1–10. doi:10.5121/ijdkp.2015.5201. Available at: https://www.researchgate.net/publication/275224157_A_Review_on_Evaluation_Metrics_for_Data_Classification_Evaluations. (Accessed: 10 April 2024)
- 10) Joshi, S. and Abdelfattah, E. (2021) 'Multi-class text classification using machine learning models for online drug reviews', 2021 IEEE World AI IoT Congress (AIoT) [Preprint]. doi:10.1109/aiiot52608.2021.9454250. Available at: <https://ieeexplore.ieee.org/abstract/document/9454250> (Accessed: 20 April 2024)
- 11) Kouiroukidis, N. and Evangelidis, G. (2011) 'The effects of dimensionality curse in high dimensional knn search', 2011 15th Panhellenic Conference on Informatics [Online]. doi:10.1109/pci.2011.45. Available at:

https://ieeexplore.ieee.org/abstract/document/6065061?casa_token=ChWutua4R64AAAAA:2aYB4x61l8a1Y5wqxv5-ME7nJ_xWsd9SToeYrIB8cgzVBMkeTCQTNaVbAdvNrgVhG1opdfC Accessed: 20 April 2024

- 12) Ligthart, A., Catal, C. and Tekinerdogan, B. (2021) 'Analyzing the effectiveness of semi-supervised learning approaches for opinion spam classification', *Applied Soft Computing*, 101, p. 107023. doi:10.1016/j.asoc.2020.107023. Available at: <https://www.sciencedirect.com/science/article/pii/S1568494620309625> Accessed: 05 April 2024
- 13) Liu, B. et al. (2017) 'A semi-supervised Convolutional Neural Network for Hyperspectral Image Classification', *Remote Sensing Letters*, 8(9), pp. 839–848. doi:10.1080/2150704x.2017.1331053. Available at: <https://www.tandfonline.com/doi/full/10.1080/2150704x.2017.1331053> Accessed: 29 April 2024
- 14) McCallum, A. and Nigam, K., 1998. A Comparison of Event Models for Naive Bayes Text Classification. AAAI-98 Workshop on Learning for Text Categorization. Available at: <http://yangli-feasibility.com/home/classes/lfd2022fall/media/aaaiws98.pdf> (Accessed: 14 April 2024)
- 15) Müller, A.C. and Guido, S. (2016) *Introduction to machine learning with python*. Nanjing: O'REILLY.
- 16) Pimpalkar, A.P. and Retna Raj, R.J. (2020) 'Influence of Pre-processing Strategies on the Performance of ML Classifiers Exploiting TF-IDF and BOW Features', *ADCAIJ: Advances in Distributed Computing and Artificial Intelligence Journal*, 9(2), pp. 49–68. doi:10.14201/adcaij.2023121. Available at: <https://www.torrossa.com/en/resources/an/5010980#> (Accessed: 07 April 2024.)
- 17) Rigatti, S.J. (2017) 'Random Forest', *Journal of Insurance Medicine*, 47(1), pp. 31–39. doi:10.17849/insm-47-01-31-39.1. Available at: <https://meridian.allenpress.com/jim/article/47/1/31/131479/Random-Forest> (Accessed: 10 April 2024)
- 18) scikit (no date) Sklearn.linear_model.logisticregression. Available at: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html (Accessed: 28 April 2024).
- 19) scikit (no date) Sklearn.linear_model.Sgdclassifier. Available at: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html (Accessed: 29 April 2024).
- 20) Singla, Z., Randhawa, S. and Jain, S. (2017) 'Sentiment analysis of customer product reviews using Machine Learning', 2017 International Conference on Intelligent Computing and Control (I2C2) [online]. doi:10.1109/i2c2.2017.8321910. Available at: <https://ieeexplore.ieee.org/abstract/document/8321910> (Accessed: 11 April 2024).
- 21) *Sklearn.naive_bayes.multinomialnb* (no date) *scikit*. Available at: https://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.MultinomialNB.html (Accessed: 24 April 2024).
- 22) Van Engelen, J.E. and Hoos, H.H., 2020. A survey on semi-supervised learning. Springer, 109(1), pp.373-440. <https://doi.org/10.1007/s10994-019-05855-6> Available at: <https://link.springer.com/article/10.1007/s10994-019-05855-6> (Accessed 20 April 2024)

- 23) Vijayarani, S., Ilamathi, J. and Nithya. (2015). Preprocessing techniques for text mining - an overview. *International Journal of Computer Science & Communication Networks*, 5(1), 7-16. ISSN: 2249-5789. Available at: https://www.researchgate.net/profile/Vijayarani-Mohan/publication/339529230_Preprocessing_Techniques_for_Text_Mining_-_An_Overview/links/5e57a0f7299bf1bdb83e7505/Preprocessing-Techniques-for-Text-Mining-An-Overview.pdf (Accessed: 11 April 2024)
- 24) Widmann, N. and Verberne, S. (2017) 'Graph-based semi-supervised learning for text classification', *Proceedings of the ACM SIGIR International Conference on Theory of Information Retrieval* [Preprint]. doi:10.1145/3121050.3121055. Available at: <https://dl.acm.org/doi/abs/10.1145/3121050.3121055> (Accessed: 04 April 2024)
- 25) Wirth, R. and Hipp, J. (2000) *CRISP-DM: Towards a Standard Process Model for Data Mining* [Online]. Available at: <https://www.cs.unibo.it/~danilo.montesi/CBD/Beatriz/10.1.1.198.5133.pdf> (Accessed: 10 April 2024).
- 26) Zhang, Z. (2016) 'Introduction to machine learning: K-Nearest Neighbors', *Annals of Translational Medicine*, 4(11), pp. 218–218. doi:10.21037/atm.2016.03.37. Available at: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4916348/> (Accessed: 10 April 2024)
- 27) Zhu, X. (2005) 'Semi-supervised Learning Literature Survey', *Computer Sciences Department*. Available at: <https://minds.wisconsin.edu/handle/1793/60444> (Accessed: 04 April 2024).

7. Python Code

```
# Load some necessary libraries
```

```
import pandas as pd
```

```
import pandas as pd
```

```
import numpy as np
```

```
import re
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
import nltk
```

```
nltk.download('punkt')
```

```
from nltk.corpus import stopwords
```

```
nltk.download('stopwords')
```

```
from nltk.stem import WordNetLemmatizer
```

```
nltk.download('wordnet')
```

```
from nltk.tokenize import word_tokenize
```

```
from nltk.stem import PorterStemmer
```

```
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
```

```
from sklearn.linear_model import SGDClassifier, LogisticRegression
```

```
from sklearn.metrics import f1_score, confusion_matrix, accuracy_score, classification_report
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.pipeline import Pipeline
```

```
from sklearn.preprocessing import FunctionTransformer
```

```
from sklearn.semi_supervised import SelfTrainingClassifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.svm import SVC
```

```
# Connect with drive
```

```
from google.colab import drive
```

```
drive.mount('/content/drive')
```

```
# load original file
```

```
path = '/content/drive/MyDrive/A_II_Emotion_Data_Student_Copy_Final.xlsx'
```

```
data1 = pd.read_excel(path)
```

```

data1.info()

# %% [markdown]
# Data preprocessing

# library to detect languages
pip install langdetect

# %%
#remove special characteres from the reviews
def clean_text1(text):
    text = re.sub(r'http\S+', '', text)
    text = re.sub(r'<[^>]+>', '', text)
    text = re.sub(r'[^a-zA-Z\s]', '', text)
    return text

# apply previous function
data1['text_reviews_1'] = data1['text_reviews_'].apply(clean_text1)

# %%
#detecte languages in the reviews
from langdetect import detect, LangDetectException

def detector_idioma(texto):
    try:
        return detect(texto)
    except LangDetectException:
        return "unknown"

data1["language"] = data1['text_reviews_1'].apply(detector_idioma)
df_no_ingles = data1[data1['language'] != 'en']

# total number of english reviews
total_en = (data1["language"] == "en").sum()

# print
total_en

```

```

# %%
#plot distribution of languages
language_counts = data1['language'].value_counts()

plt.figure(figsize=(4, 2))
plt.bar(language_counts.index, language_counts.values, color='skyblue', edgecolor='black')
plt.xlabel('Language')
plt.ylabel('Number of Reviews')
plt.title('Language Distribution Accross the reviews')
plt.xticks(rotation=90, fontweight='bold')
plt.show()

# %%
#delete from the data set with no english reviews those reviews from GB to reduce the size of it
df_no_ingles1 = df_no_ingles[df_no_ingles['country_'] != 'GB']

# %%
#calculate the lenght of the reviews
df_no_ingles1['longitud'] = df_no_ingles1['text_reviews_1'].str.len()

total_caracteres = df_no_ingles1['longitud'].sum()

# %%
total_caracteres

# %%
# API deepl only allows 500000 for free, so the data set is divide and Baidu is also used
df_no_ingles1['length'] = df_no_ingles1['text_reviews_1'].str.len()
df_no_ingles1['cumulative_length'] = df_no_ingles1['length'].cumsum()

# Filter the dataset
df_filtered = df_no_ingles1[df_no_ingles1['cumulative_length'] <= 500000]
df_filtered2 = df_no_ingles1[df_no_ingles1['cumulative_length'] > 500000]

#Translate the others 400000 caracteres
import requests
from hashlib import md5
import random
import json
appid = '234509250023333'

```

```
appkey = 'BRdzUuBI3BBAMC9es1d8'
```

```
def make_md5(s, encoding='utf-8'):
    return md5(s.encode(encoding)).hexdigest()
```

```
def baidu_api(query, from_lang='auto', to_lang='en'):
    endpoint = 'http://api.fanyi.baidu.com'
    path = '/api/trans/vip/translate'
    url = f'{endpoint}{path}'
    salt = random.randint(32768, 65536)
    sign = make_md5(appid + query + str(salt) + appkey)
    headers = {'Content-Type': 'application/x-www-form-urlencoded'}
    payload = {
        'appid': appid,
        'q': query,
        'from': from_lang,
        'to': to_lang,
        'salt': salt,
        'sign': sign
    }

    try:
        r = requests.post(url, params=payload, headers=headers, timeout=10)
        result = r.json()
        if 'error_code' in result:
            print(f"Error {result['error_code']}: {result.get('error_msg', 'No error message')}")
            return query
        return result["trans_result"][0]['dst']
    except Exception as e:
        print(f"Error during translation: {e}")
        return query
```

```
def translate_text(text):
    return baidu_api(text, from_lang='auto', to_lang='en')
```

```
df_filtered2['translated_text'] = df_filtered2['text_reviews_1'].apply(translate_text)
```

```
# %%
```

```
#Translate the first 500000 characteres
```

```

import deepL

translator = deepL.Translator("ee87c8a8-fe26-4956-a1d6-bbe41f8590f1:fx")

def translate_text(text):
    try:
        result = translator.translate_text(text, target_lang="EN-GB")
        return result.text
    except Exception as e:
        print(f"Error al traducir el texto: {e}")
        return text

df_filtered["translated_text"] = df_filtered["text_reviews_1"].apply(translate_text)

# %%
#adding the Deepl translated reviews to the original data set
result_df = pd.merge(data, df_filtered[["ID_", "translated_text"]], on='ID_', how='left')

result_df["text_reviews_"] = result_df["translated_text"].combine_first(result_df["text_reviews_"])

result_df.drop(["translated_text", "text_reviews_1", "language"], axis=1, inplace=True)

# %%
#adding the Baidu translated reviews to the previous data set
merged_df = pd.merge(result_df, df_filtered2[["ID_", "translated_text"]], on='ID_', how='left')

merged_df["text_reviews_"] = merged_df["translated_text"].combine_first(merged_df["text_reviews_"])

final_df = merged_df.drop("translated_text", axis=1)

# %%
final_df.drop("Cleaned_reviews", axis=1, inplace=True)

# this data set contains same information as the original provided but with all the reviews in English
final_df.to_csv('/content/drive/MyDrive/final_reviews.csv', index=False)

# %% [markdown]
# New file with all the reviews in English

```



```

# load
path = '/content/drive/MyDrive/final_reviews.csv'
data = pd.read_csv(path)

data.info()

#remove 2 missing values originated when the reviews were translated
data = data.dropna(subset=['text_reviews_'])

# %%
#plot language distribution accross reviews
from langdetect import detect, LangDetectException

def detector_idioma(texto):
    try:
        return detect(texto)
    except LangDetectException:
        return "unknown"

to_plot = data.copy()

to_plot["language"] = to_plot["text_reviews_"].apply(detector_idioma)

language_counts = to_plot["language"].value_counts()

plt.figure(figsize=(6, 4))
plt.bar(language_counts.index, language_counts.values, color='skyblue', edgecolor='black')
plt.xlabel('Language')
plt.ylabel('Number of Reviews')
plt.title('Language Distribution Accross the translated reviews')
plt.xticks(rotation=90)
plt.show()

## Text Data Pre-Processing

# function to do text pre-processing
def clean_text(text):

```

```

text = re.sub(r'http\S+', '', text) # Remove URLs
text = re.sub(r'<[^>]+>', '', text) # Remove HTML tags
text = text.lower() # Lowercase text
text = re.sub(r'\b\w{1,2}\b', '', text) # Remove words with 1 or 2 letters
text = re.sub(r'[^a-z\s]', '', text) # Keep text with letters and spaces

```

```

# Tokenize

```

```

tokens = word_tokenize(text)

```

```

# Remove stopwords

```

```

stop_words = set(stopwords.words('english'))
tokens = [word for word in tokens if word not in stop_words]

```

```

# Lemmatize 0.56 accuracy

```

```

lemmatizer = WordNetLemmatizer()
tokens = [lemmatizer.lemmatize(word) for word in tokens]

```

```

# Stemming 0.54 accuracy

```

```

#stemmer = PorterStemmer()
#tokens = [stemmer.stem(word) for word in tokens]

```

```

return ' '.join(tokens)

```

```

# new column is created with the clean text

```

```

data['Cleaned_reviews'] = data['text_reviews_'].apply(clean_text)

```

```

## Initial descriptive analysis

```

```

sns.set(style="whitegrid")

```

```

plt.figure(figsize=(8, 8))

```

```

top_countries = data['country_'].value_counts().head(15).index

```

```

# Filtering the data to include only the top 15 countries

```

```

filtered_data = data[data['country_'].isin(top_countries)]

```

```

missing_emotions = data['emotions_'].isnull().sum()

```

```

# Adjusting the categorical columns to use the filtered data for country_

```

```

categorical_columns = ['brand_name_', 'country_', 'emotions_']

```

```

numerical_columns = ['star_rating_']

```

#4 subplots

rows, cols = 2, 2

#loop to iterate in the variables

for index, column in enumerate(categorical_columns):

plt.subplot(rows, cols, index + 1)

#take country from the data set where there are only top 15 countries

if column == 'country_':

sns.countplot(y=column, data=filtered_data, palette="Set2")

#take the rest of the variables from the original dataset

else:

count_plot = sns.countplot(y=column, data=data, palette="Set2")

plt.title(f'{column}', fontweight='bold')

plt.tight_layout()

#add the emotions to the missing values in the bottom left-hand corner

if column == 'emotions_':

plt.text(count_plot.get_xlim()[1], count_plot.get_ylim()[0], f'Missing Values: {missing_emotions}',

horizontalalignment='left', verticalalignment='top',

fontsize=12, color='red')

#plot

plt.subplot(rows, cols, len(categorical_columns) + 1)

sns.violinplot(y='star_rating_', data=data, color="skyblue")

plt.title('star_rating_', fontweight='bold')

plt.tight_layout()

plt.show()

Machine Learning

Unlabeled data

data frame is created with missing observations of emotions (Unlabeled data) and Cleaned reviews is included in the data frame.

unlabeled_data = data[pd.isna(data['emotions_'])][['Cleaned_reviews']]

unlabeled_data['emotions_'] = -1

Labeled Data

```

# %%
# data frame is created with known observations of emotions (Labeled data) and Cleaned reviews is included in
the data frame.
labeled_data = data[data['emotions_'].notna()][['Cleaned_reviews', 'emotions_']]

# Create X and y
X = data['Cleaned_reviews']
y = data['emotions_']

# known emotions
y_labeled = labeled_data['emotions_']
# missing emotions
y_unlabeled = unlabeled_data['emotions_']
# cleaned reviews for known emotions
X_labeled = labeled_data['Cleaned_reviews']
# cleaned reviews for missing emotions
X_unlabeled = unlabeled_data['Cleaned_reviews']

# Function for classification report
def eval_and_print_metrics(clf, X_train, y_train, X_test, y_test):
    print("Number of training samples:", len(X_train))
    print("Unlabeled samples in training set:", sum(1 for x in y_train if x == -1)) #if x == 'NaN'
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)

    print(
        "Micro-averaged F1 score on test set: %0.3f"
        % f1_score(y_test, y_pred, average="micro")
    )
    print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
    print("\nClassification Report:\n", classification_report(y_test, y_pred, zero_division=1))
    print("\n\n")

# Split the labeled data
X_train, X_test, y_train, y_test = train_test_split(X_labeled, y_labeled, test_size=0.2, stratify=y_labeled,
random_state=42)

from sklearn.feature_extraction.text import TfidfVectorizer

```

```

from sklearn.pipeline import make_pipeline
from sklearn.model_selection import GridSearchCV

# set TfidfVectorizer parameter using a logistic regression

pip install mglearn

import mglearn

pipe1 = make_pipeline(TfidfVectorizer(), LogisticRegression(max_iter=10000))
#Source: (Bird, S., Klein, E. and Loper, E. (2009) Natural language processing with python. Beijing: O'Reilly.)

#set parameter of logistic regression (C), based on previous calculation it is known that we required a high value
for C
param_grid1 = {
    "logisticregression__C": [10, 100],
    "tfidfvectorizer__ngram_range": [(1, 1), (1, 2), (1, 3)],
    "tfidfvectorizer__max_features": [1000, 2000, 3000, 5000]
}

grid1 = GridSearchCV(pipe1, param_grid1, cv=5)
grid1.fit(X_train, y_train)

scores = grid1.cv_results_['mean_test_score'].reshape(len(param_grid1["logisticregression__C"]),
                                                    len(param_grid1["tfidfvectorizer__ngram_range"]),
                                                    len(param_grid1["tfidfvectorizer__max_features"]))
scores_mean = scores.mean(axis=0)

heatmap = mglearn.tools.heatmap(scores_mean, xlabel="max_features", ylabel="ngram_range", cmap="viridis",
                                fmt="%0.3f",
                                xticklabels=param_grid1["tfidfvectorizer__max_features"],
                                yticklabels=param_grid1["tfidfvectorizer__ngram_range"])

plt.colorbar(heatmap)
plt.show()

# best ngram (1,2)
# max_features 3000
#

```

```
## Supervised Machine learning** model for labeled data
```

```
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
```

```
from sklearn.svm import SVC
```

```
from sklearn.naive_bayes import MultinomialNB
```

```
vectorizer_params = {  
    'ngram_range': (1, 2),  
    'max_features': 3000  
}
```

```
#hyper parameter tuning
```

```
param_grid = {  
    "Random Forest": {  
        "clf__n_estimators": [50, 100, 200],  
        "clf__max_depth": [None, 10, 20, 30],  
        "clf__min_samples_split": [2, 5, 10]  
    },  
    "K-Nearest Neighbors": {  
        "clf__n_neighbors": [3, 5, 6, 7],  
        "clf__weights": ["uniform", "distance"],  
        "clf__algorithm": ["auto", "ball_tree"]  
    },  
    "Support Vector Machine": {  
        "clf__C": [0.1, 1, 10],  
        "clf__kernel": ["linear", "rbf", "poly"]  
    },  
    "SDG": {  
        "clf__alpha": [1e-5, 1e-4, 1e-3],  
        "clf__penalty": ["l2", "l1"],  
        "clf__loss": ["hinge", "log"]  
    },  
    "Naive Bayes": {  
        "clf__alpha": [0.01, 0.1, 1.0]  
    }  
}
```

#classifiers

```
supervised_classifiers = [  
    ("Random Forest", RandomForestClassifier(), param_grid["Random Forest"]),  
    ("K-Nearest Neighbors", KNeighborsClassifier(), param_grid["K-Nearest Neighbors"]),  
    ("Support Vector Machine", SVC(), param_grid["Support Vector Machine"]),  
    ("SDG", SGDClassifier(), param_grid["SDG"]),  
    ("Naive Bayes", MultinomialNB(), param_grid["Naive Bayes"])  
]
```

classifier_names = []

accuracy_scores = []

#loop through the classifiers and parameters, saving best accuracies and names in the lists created above

for name, clf, params in supervised_classifiers:

```
    pipeline = Pipeline([  
        ("vect", TfidfVectorizer(**vectorizer_params)),  
        ("clf", clf),  
    ])
```

```
    grid_search = GridSearchCV(pipeline, param_grid=params, n_jobs=-1, cv=10)
```

```
    grid_search.fit(X_labeled, y_labeled)
```

```
    best_clf = grid_search.best_estimator_
```

```
    best_params = grid_search.best_params_
```

```
    best_score = grid_search.best_score_
```

Print best hyperparameters

```
print(f"Best Hyperparameters for {name}: {best_params}")
```

Evaluate and print metrics

```
eval_and_print_metrics(best_clf, X_train, y_train, X_test, y_test)
```

Store the name and the accuracy for plotting

```
classifier_names.append(name)
```

```
accuracy_scores.append(best_score)
```

change classifiers names for graph clarity

```
classifier_names = ["RF", "K-NN", "SVM", "SDG", "MNB"]
```

```
# Plotting the accuracies
```

```
plt.figure(figsize=(4, 3))  
plt.bar(classifier_names, accuracy_scores, color='blue')  
plt.title('Cross-validation Classifiers Accuracies')  
plt.xticks(rotation=90)  
plt.tight_layout()  
plt.show()
```

```
# build model with the best model using the best hyper parameters
```

```
best_params_sdg = { 'alpha': 0.001, 'loss': 'log',  
                    'penalty': 'l1' }
```

```
pipeline_sdg = Pipeline([  
    ("vect", TfidfVectorizer(**vectorizer_params)),  
    ("clf", SGDClassifier(**best_params_sdg))  
])
```

```
pipeline_sdg.fit(X_train, y_train)
```

```
y_pred1 = pipeline_sdg.predict(X_test)
```

```
#cm
```

```
cm = confusion_matrix(y_test, y_pred1)  
plt.figure(figsize=(4,4))  
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", cbar=False)  
plt.title("Confusion Matrix")  
plt.xlabel("Predicted Label")  
plt.ylabel("True Label")  
plt.show()
```

```
##Semi-supervised classifiers
```

```
#mantain x tests
```

```
test_indices = X_test.index
```

```
# Exclude test data from X_labeled and y_labeled based on the identified indices
```

```
X_labeled_filtered = X_labeled.drop(index=test_indices, errors='ignore')  
y_labeled_filtered = y_labeled.drop(index=test_indices, errors='ignore')
```



```

# Concatenate the filtered labeled data with the unlabeled data
# lo que se esta haciendo es extraer de X labeled e y labeled el test data test (me quedo con el train)
# se junta train labeled con toda la unlabeled por lo tanto tengo: X = x train label + all unlabeled. Y = y train
labeled + all unlabeled

# x_test y y_test es todo labeled data para comprobar accuracy de las predicciones
X=X_combined = pd.concat([X_labeled_filtered, X_unlabeled])
y=y_combined = pd.concat([y_labeled_filtered, y_unlabeled])

# Define the mapping for labels
label_mapping = {
    'anger': 0,
    'disgust': 1,
    'fear': 2,
    'joy': 3,
    'neutral': 4,
    'sadness': 5,
    'surprise': 6, -1:-1}

# Apply the mapping to labels
y = [label_mapping[label] for label in y]
print(y)
y_test = [label_mapping[label] for label in y_test]
print(y_test)

#selflearning
from sklearn.linear_model import LogisticRegression, SGDClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.semi_supervised import SelfTrainingClassifier

vectorizer_params = {
    'ngram_range': (1, 2),
    'max_features': 3000
}

random_state_seed = 40425150

# Defining classifiers with the best hyperparameter of supervised

```

```

non_graph_based_classifiers = [
    ("SelfTraining Logistic Regression", SelfTrainingClassifier(LogisticRegression(max_iter=1000),
criterion='threshold', threshold=0.9)),
    ("SelfTraining K-Nearest Neighbors",
SelfTrainingClassifier(base_estimator=KNeighborsClassifier(n_neighbors=7, algorithm='auto',
weights='distance'))),
    ("SelfTraining Naive Bayes", SelfTrainingClassifier(MultinomialNB())),
    ("SelfTraining SVC linear", SelfTrainingClassifier(SVC(probability=True, C=10, gamma='scale', kernel='linear'),
criterion='threshold', threshold=0.9)),
    ("SelfTraining SDG", SelfTrainingClassifier(SGDClassifier(alpha=0.001, loss='log', penalty='l1', verbose =
True))),
    ("Random Forest", SelfTrainingClassifier(RandomForestClassifier(max_depth=20, min_samples_split=10,
n_estimators=200)))
]

```

#lop and results of self-learning classifiers

```

for name, clf in non_graph_based_classifiers:
    print(f"Evaluating {name}")
    st_pipeline = Pipeline([
        ("vect", TfidfVectorizer(**vectorizer_params)),
        ("clf", clf),
    ])
    eval_and_print_metrics(st_pipeline, X, y, X_test, y_test)

```

Label spreading

```

from sklearn.semi_supervised import LabelSpreading

```

```

vectorizer_params = {
    'ngram_range': (1, 2),
    'max_features': 3000
}

```

```

ls_pipeline = Pipeline(
    [
        ("vect", TfidfVectorizer(**vectorizer_params)),
        ("toarray", FunctionTransformer(lambda x: x.toarray())),
        ("clf", LabelSpreading()),
    ]
)

```

```
ls_pipeline.fit(X, y)
```

```
eval_and_print_metrics(ls_pipeline, X, y, X_test, y_test)
```

```
# Best performance is SDG
```

```
# Report of Self learning SDG
```

```
st_svc_pipeline = Pipeline([  
    ("vect", TfidfVectorizer(**vectorizer_params)),  
    ("SelfTraining SGD", SelfTrainingClassifier(SGDClassifier(alpha=0.001, loss='log', penalty='l1', verbose =  
    True)))])
```

```
st_svc_pipeline.fit(X, y)
```

```
y_pred2 = st_svc_pipeline.predict(X_test)
```

```
print("Classification Report:")  
print(classification_report(y_test, y_pred2))
```

```
cm2 = confusion_matrix(y_test, y_pred2)  
plt.figure(figsize=(4,4))  
sns.heatmap(cm2, annot=True, fmt="d", cmap="Blues", cbar=False)  
plt.title("Confusion Matrix")  
plt.xlabel("Predicted Label")  
plt.ylabel("True Label")  
plt.show()
```

```
# Predicting with SDG the unlabeled data
```

```
predictions = st_svc_pipeline.predict(X_unlabeled)
```

```
plt.hist(predictions, bins='auto', color='skyblue', edgecolor='black')  
plt.title("Distribution of Predictions")  
plt.xlabel("Predicted Label")  
plt.ylabel("Frequency")  
plt.show()
```

```
# inverse mapping
```

```
data1 = data.copy()
```

```
inverse_label_mapping = {v: k for k, v in label_mapping.items()}
```

```
text_predictions = np.array([inverse_label_mapping[pred] for pred in predictions])
```

```
missing_emotion_indices = data1[data1['emotions_'].isna()].index
```

```
assert len(predictions) == len(missing_emotion_indices)
```

```
data1.loc[missing_emotion_indices, 'emotions_'] = text_predictions
```

```
# plot distribution of emotions after predictions
```

```
plt.hist(data1['emotions_'], bins='auto', color='skyblue', edgecolor='black')
```

```
figsize=(7, 7)
```

```
plt.title("Distribution of Predictions")
```

```
plt.xlabel("Predicted Label")
```

```
plt.ylabel("Frequency")
```

```
plt.xticks(fontweight = "bold")
```

```
plt.show()
```

```
# Analysis for each BRAND
```

```
sns.set(style="whitegrid")
```

```
# Emotions for each brand
```

```
plt.figure(figsize=(4, 3))
```

```
sns.countplot(y='emotions_', hue='brand_name_', data=data1, palette='Set2')
```

```
plt.title("Distribution of emotions by brand")
```

```
plt.xlabel("Frequency")
```

```
plt.ylabel("Emotions")
```

```
plt.legend(title='Brand')
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# percentages
```

```
total_emotions_per_brand = data1.groupby('brand_name_').size().reset_index(name='total_per_brand')
```

```
# Count each emotion per brand
```

```
emotion_counts = data1.groupby(['brand_name_', 'emotions_']).size().reset_index(name='counts')
```

```
# Merge the total counts into emotion counts to calculate percentages
```

```
emotion_counts = emotion_counts.merge(total_emotions_per_brand, on='brand_name_')
```

```
# Calculate percentage
```

```
emotion_counts['percentage'] = 100 * emotion_counts['counts'] / emotion_counts['total_per_brand']
```

```
sns.set(style="whitegrid")
```

```
# Create a bar plot to visualize percentages
```

```
plt.figure(figsize=(4, 3))
```

```
sns.barplot(x='percentage', y='emotions_', hue='brand_name_', data=emotion_counts, palette='Set2')
```

```
plt.title('Percentage Distribution of Emotions by Brand')
```

```
plt.xlabel('Percentage')
```

```
plt.ylabel('Emotions')
```

```
plt.legend(title='Brand')
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# to identify the 2 brands
```

```
brands = data1['brand_name_'].unique()
```

```
# loop to iterate based on the 2 brands availables to plot star rating for each brand for each emotion
```

```
for brand in brands:
```

```
    # Filter the dataset for the current brand
```

```
    brand_data = data1[data1['brand_name_'] == brand]
```

```
    plt.figure(figsize=(6, 4))
```

```
    sns.violinplot(x='star_rating_', y='emotions_', data=brand_data, palette='Set3')
```

```
    plt.title(f'Star rating by emotion for {brand}\n')
```

```
    plt.xlabel('Star Rating')
```

```
    plt.ylabel('Emotion')
```

```
    plt.tight_layout()
```

```
    plt.show()
```

LOGISTIC REGRESSION TO ANALYSE COEFFICIENTS

#Idea and part of the code is taking from: (Müller, A.C. and Guido, S. (2016) Introduction to machine learning with python. Nanjing: O'REILLY)

```
def visualize_coefficients_multiclass_modified(coef_matrix, feature_names, class_labels, n_top_features=20):
    n_classes = coef_matrix.shape[0]
    for i in range(n_classes):
        plt.figure(figsize=(8, 4))
        coef = coef_matrix[i, :]
        mglearn.tools.visualize_coefficients(coef, feature_names, n_top_features=n_top_features)
        plt.title(f'Emotion: {class_labels[i]}', fontsize=14, fontweight='bold')
        plt.xticks(fontsize=14, fontweight='bold')
        plt.show()
```

Loop through each brand to train and evaluate a model

```
for brand in brands:
    # Filter the dataset for the current brand
    brand_data = data1[data1['brand_name_'] == brand]

    # Prepare the data
    X_train2, X_test2, y_train2, y_test2 = train_test_split(
        brand_data['Cleaned_reviews'], brand_data['emotions_'], test_size=0.2, random_state=42
    )

    # Create a pipeline with TF-IDF and Logistic Regression
    model = make_pipeline(TfidfVectorizer(**vectorizer_params), LogisticRegression(max_iter=10000, C=10))

    # Train the model
    model.fit(X_train2, y_train2)

    # Evaluate the model
    predictions = model.predict(X_test2)
    print(f'Classification Report for Brand: {brand}\n')
    print(classification_report(y_test2, predictions))

    # Visualization of coefficients
    vect = model.named_steps['tfidfvectorizer']
```


```





features_names = np.array(vect.get_feature_names_out())
coef = model.named_steps['logisticregression'].coef_
class_labels = model.named_steps['logisticregression'].classes_
visualize_coefficients_multiclass_modified(coef, features_names, class_labels, n_top_features=20)


```

8. Appendix

-. Brand Z_:







-. Brand H_:

